



Project:

big-data-infrastructure-demo



Purpose

This repository is designed to **demonstrate the skills and concepts** required for the role of a **Big Data Infrastructure Engineer** — using a realistic, modular, and containerized setup. It simulates a **real-time health data pipeline**, transforming clinical records into queryable datasets using tools from the modern data engineering stack.



Environment Description

This project mimics a healthcare environment where patient records from an **OpenMRS database** (via MySQL) are captured in real time and streamed through a **big data infrastructure** for storage, querying, and analysis.

Key goals:

- Show proficiency with **Hadoop ecosystem components** (HDFS, Hive, YARN)
 - Demonstrate **Kafka-based real-time ingestion**
 - Use **Debezium** for CDC (Change Data Capture)
 - Enable **SQL querying on big data** using Hive
 - Optional: Integrate **Prometheus and Grafana** for system monitoring
 - Containerized setup using **Docker Compose**
 - Deployable and testable locally or in the cloud (e.g., DigitalOcean, AWS)
-



Repository Structure

Final System Architecture

SCSS

CopyEdit

MySQL (OpenMRS DB Dump)

↓ (via Debezium)

Kafka ↔ Kafka Connect → HDFS (Hadoop)

↓

Hive

↓

Queries via Beeline

↓

(Optional) Monitoring via Prometheus + Grafana

Phase-by-Phase Development Approach

✓ Phase 1: Minimal Pipeline — MySQL + Debezium + Kafka

- MySQL container with `openmrs_sample_dump.sql`
- Debezium running inside Kafka Connect container
- Kafka and Zookeeper for messaging
- Insert into DB → See CDC event in Kafka

 File: `compose/stage-1-minimal.yaml`

✓ Phase 2: Add Hadoop — Stream into HDFS

- Deploy Hadoop (NameNode + 2 DataNodes)

- Kafka Connect uses HDFS Sink connector
- JSON/Avro events written to HDFS

 File: `compose/stage-2-hdfs.yaml`

✓ Phase 3: Add Hive — Query via SQL

- Launch Hive Metastore + HiveServer2
- Create Hive external table over HDFS data
- Use Beeline or JDBC to query patient data

 File: `compose/stage-3-hive.yaml`

✓ Phase 4: Monitoring (Optional)

- Prometheus scrapes metrics from Kafka, Hadoop, Debezium
- Grafana visualizes ingestion rate, disk usage, job status
- JMX exporters or Prometheus exporters installed


 File: `compose/stage-4-monitoring.yaml`

✓ Phase 5: GitHub CI/CD & Final Composition

- GitHub Actions: validate docker-compose + lint configs
- All phases documented and runnable
- Easy-to-clone, run, and test anywhere

 File: `compose/stage-final.yaml`
 CI: `.github/workflows/ci.yaml`

Demo: How to Show the Setup Works

 This is your **interview/demo script** — a live proof that the pipeline works.

1. Run the Full Stack

bash
CopyEdit
`docker-compose -f compose/stage-final.yaml up -d`

2. Import OpenMRS Sample DB into MySQL

bash
CopyEdit
`docker cp ./data/openmrs_sample_dump.sql mysql:/openmrs.sql`

`docker exec -i mysql sh -c 'exec mysql -u root -p$MYSQL_ROOT_PASSWORD openmrs' < ./data/openmrs_sample_dump.sql`

3. Manually Insert a Patient Record

bash
CopyEdit
`docker exec -it mysql mysql -u root -popenmrs`

`USE openmrs;`

`INSERT INTO patient (patient_id, gender, birthdate, creator, date_created)
VALUES (90001, 'F', '1987-05-12', 1, NOW());`

`INSERT INTO person_name (person_name_id, person_id, given_name, family_name, creator, date_created)`

```
VALUES (80001, 90001, 'Amina', 'Tshisekedi', 1, NOW());
```

4. Verify Kafka Received the Event

bash

CopyEdit

```
docker exec -it kafka-broker kafka-console-consumer \  
  --bootstrap-server localhost:9092 \  
  --topic dbserver1.openmrs.patient \  
  --from-beginning
```

✓ You should see a JSON message containing the new patient.

5. Confirm Data Landed in HDFS

bash

CopyEdit

```
docker exec -it hadoop-namenode hdfs dfs -ls /kafka/openmrs.patient/
```

✓ .json or .avro files appear in HDFS based on Kafka sink connector.

6. Query Data in Hive

bash

CopyEdit

```
docker exec -it hive-server beeline -u jdbc:hive2://localhost:10000
```

-- Inside Beeline:

```
SELECT payload.patient_id, payload.gender FROM patients;
```

✓ You should see the inserted patient info.

7. (Optional) Visualize in Grafana

- Open: <http://localhost:3000>
- Default login: [admin](#) / [admin](#)
- Explore dashboards: Kafka, Hadoop, Connect